



EE 451 Introduction to Parallel and Distributed Computation

Discussion 02/26/2021
University of Southern California

Map-Reduce Background



- Large set of data needs to be processed in a fast and efficient way
- In order to process large set of data in a reasonable amount time, this needs to be distributed across thousands of machines
- Programmers need to focus on solving problems without worrying about the implementation



MapReduce

- Programming Abstraction
- Two operations
 - Map
 - Reduce



MapReduce

- Map
 - Input : key-value pairs
 - Output: intermediate key-value pairs
- MapReduce framework groups all pairs with same key
- Reduce
 - Input: key, iterator values (list of values)
 - Output: list with results



MapReduce Example

- Counting each word in a large set of documents

map (String key, String value)

// key: document name

// value: document contents

for each word **w** in value:

EmitIntermediate (**w**, "1")

reduce (String key, Iterator values)

// key: word

// value: a list of counts

for each **v** in values:

result + = ParseInt(v);

Emit(AsString(result));



MapReduce Example

- Counting each word in a large set of documents

Document_1	Document_2
foo	test
bar	foo
baz	baz
foo	bar
bar	foo
test	

Expected results:

`<foo, 4>, <bar, 3>, <baz, 2>, <test, 2>`



MapReduce Example

- Counting each word in a large set of documents

```
map(String key, String value):  
    // key: document name  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");
```

```
Map(document_1,contents(document_1))
```

```
<foo, "1">  
<bar, "1">  
<baz, "1" >  
<foo, "1">  
<bar, "1">  
<test, "1">
```

```
Map(document_2,contents(document_2))
```

```
<test, "1">  
<foo, "1">  
<baz, "1">  
<bar, "1">  
<foo, "1">
```



MapReduce Example

- Counting each word in a large set of documents

```
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```

Reduce(word, values)

```
<foo, "2">  
<bar, "2">  
<baz, "1" >  
<test, "1">
```

Reduce(word, values)

```
<test, "1">  
<foo, "2">  
<baz, "1">  
<bar, "1">
```




MapReduce Example

- Counting each word in a large set of documents

```
<foo, "2">  
<bar, "2">  
<baz, "1">  
<test, "1">
```

```
<test, "1">  
<foo, "2">  
<baz, "1">  
<bar, "1">
```

Reduce(word, values)

```
<foo, "4">  
<bar, "3">  
<baz, "2">  
<test, "2">
```

Expected results:

```
<foo, 4>, <bar, 3>, <baz, 2>, <test, 2>
```



Benefit of MapReduce

- Easy to use for programmers that do not need to worry about the details of distributed computing
- Flexible and scalable in large clusters of machines



Apache Spark (1)

- Open source cluster computing framework
- Provides an interface for programming entire cluster with implicit data parallelism and fault tolerance



Apache Spark (2)

- Spark Components
 - Core: distributed task dispatching, scheduling and basic I/O functionalities
 - Spark SQL: provides support for structured and semi-structured data
 - Spark streaming: provides support for streaming analytics
 - Mlib: Distributed machine learning framework
 - GraphX: Distributed graph processing framework

Resilient Distributed Datasets (RDD) (1)



- Fault tolerant read-only collection of datasets that can be operated in parallel
- Creating RDDs
 - Transformation from an existing RDD
 - Referencing an external dataset (filesystem, HDFS)
- Operating on RDD
 - Transformation: creates new dataset from existing ones
 - Action: returns a value after running a computation on a dataset

Resilient Distributed Datasets (RDD) (2)

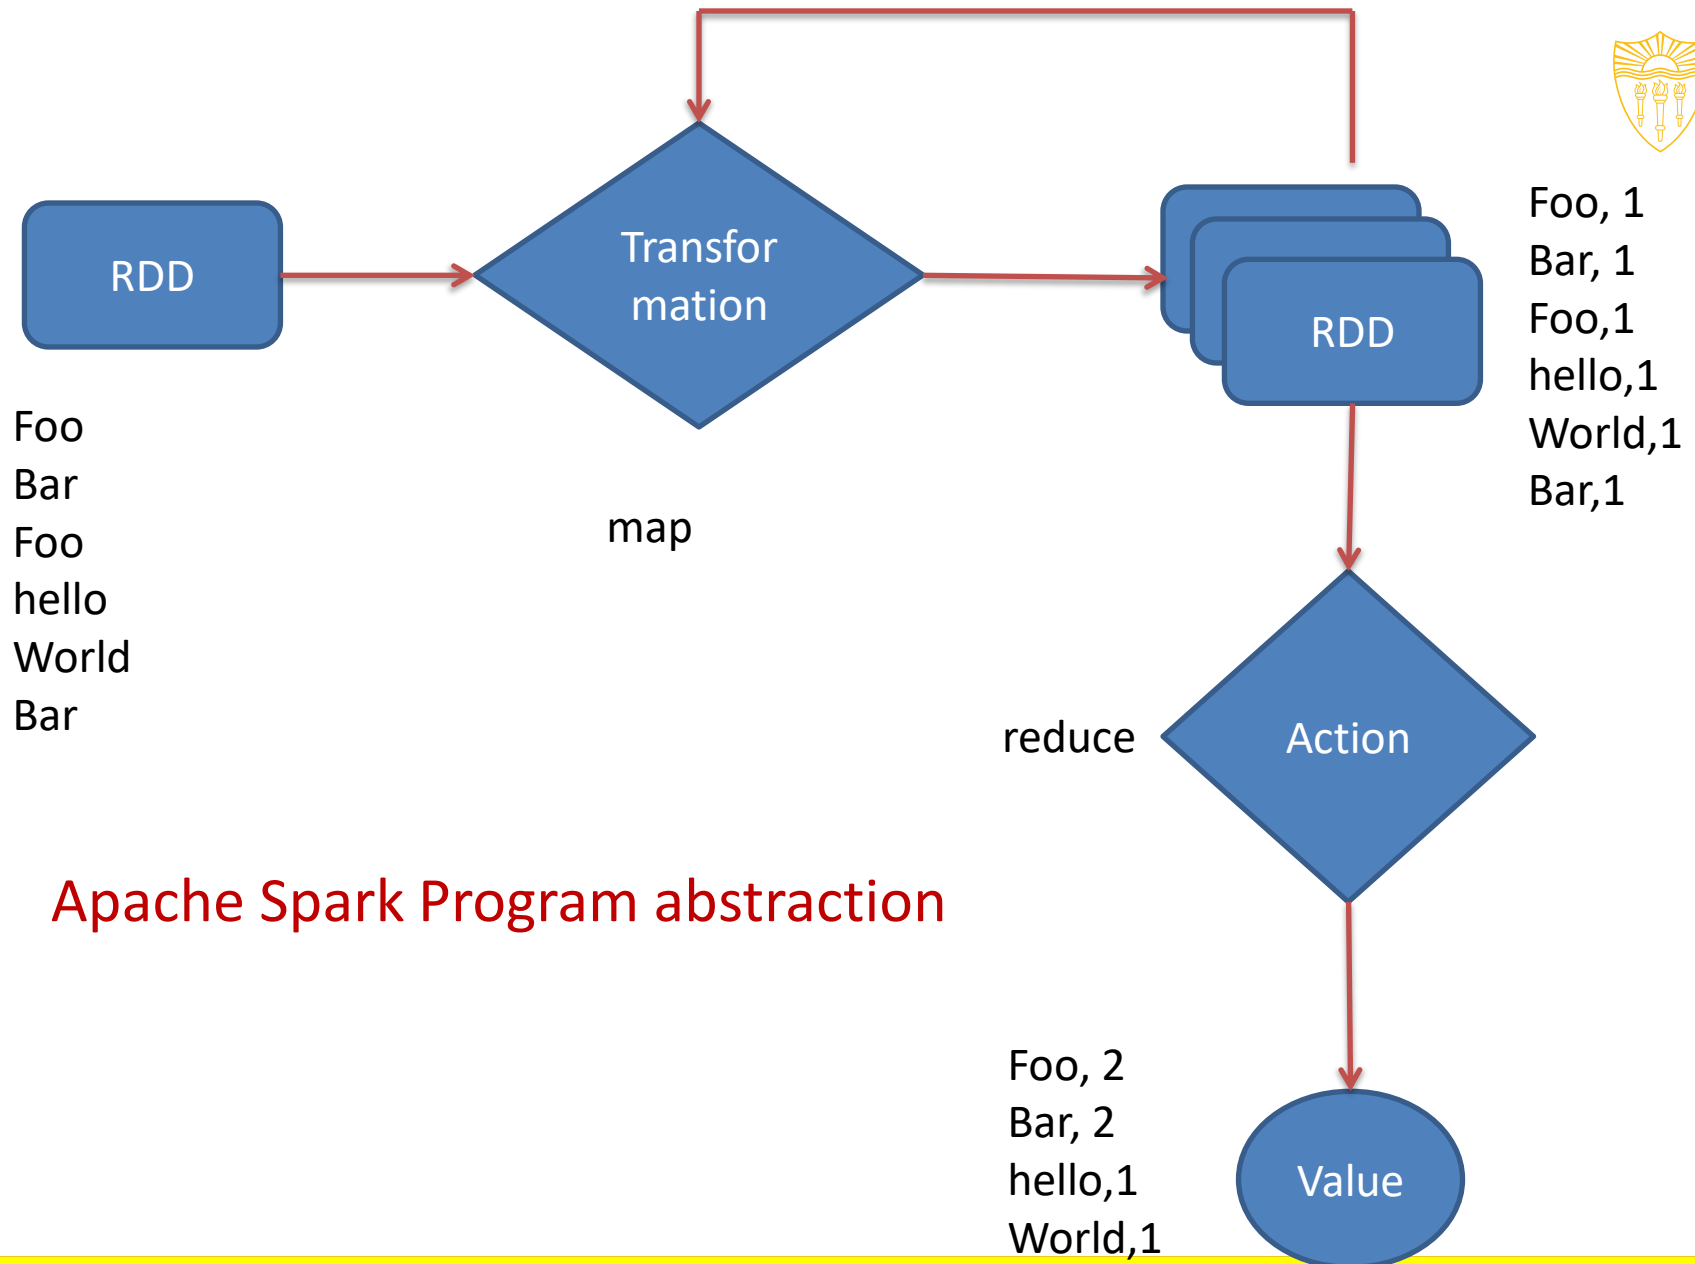


- Transformations:
 - Map: each element passed through a function.
 - Union: union of elements in source RDD
 - Intersection: intersection of elements in source RDD
 - Filter: elements which match a criteria in source RDD
- Actions
 - Reduce: aggregate elements of RDD using a function
 - Collect: create an array out of RDD
 - Count : count the number of elements in RDD

Resilient Distributed Datasets (RDD) (3)



- Transformations are lazy. Applied only when an action requires its results
- Transformation is recomputed each time an action is run on it. The results can be persisted in memory using functions such as: **cache()** or **persist()**



Apache Spark Program abstraction



Running a job

Submit to spark cluster

Cluster is the local machine

- `../bin/spark-submit --master local[*]
kmeans.py data.txt centroid.txt`

Python file

arguments




PHW #5



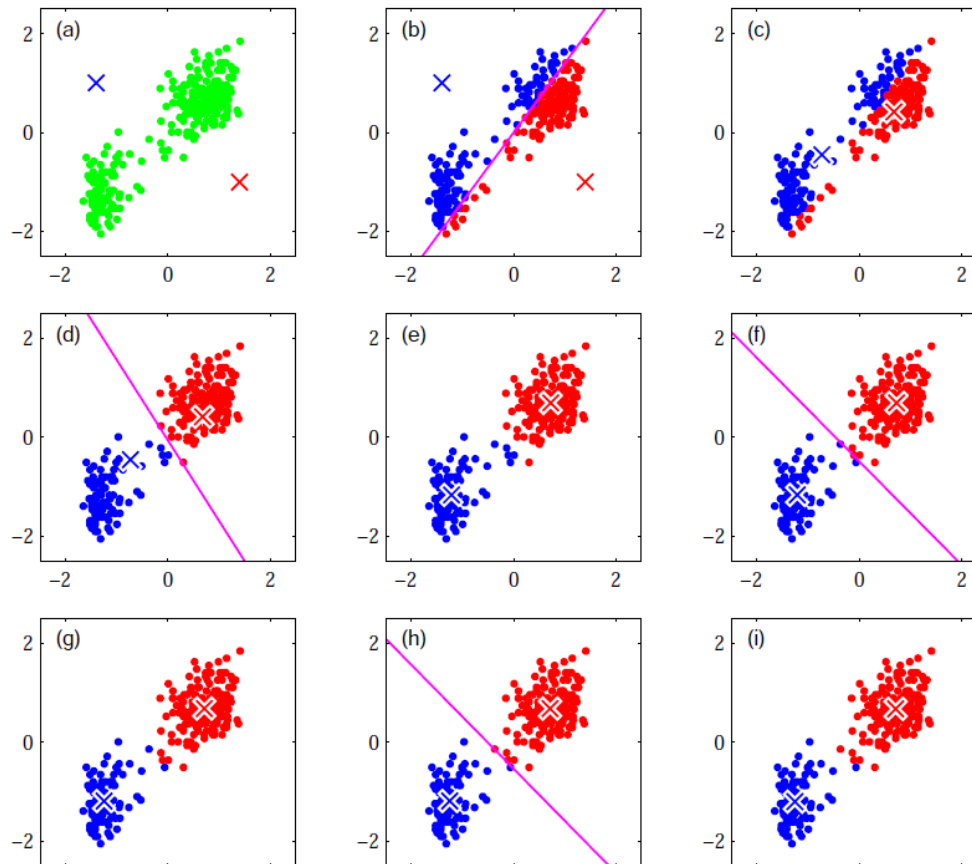
K-means Clustering (1)

- Input: a set of observations (values) $X=\{x_0, \dots, x_{N-1}\}$
- Objective: partition observations into K clusters
 - Each cluster has a mean value, μ_i ($0 < i < K$)
 - Each observation belongs to the cluster with the closest mean





K-means Clustering (2)





K-means Clustering (3)

- **Map** ($x, \mu_0, \dots, \mu_{K-1}$)

Distance = ∞

For $j = 0$ to $K - 1$ do

 If Distance < $|x - \mu_j|$

 Distance = $|x - \mu_j|$

 Cluster ID $key = j$

 End if

End for

Output key-value pair (key, x)

End for

Key

Value




K-means Clustering (4)

- **Shuffle:** $L_i \leftarrow x \mid (i, x)$
- **Reduce** (L_0, \dots, L_{K-1})
 - For $i = 0$ to $K - 1$ do
 - μ_i = average of elements in L_i
 - End for
- **ReduceByKey(L)**
 - return average of element in L



K-means Clustering (5)

- Example
 - $X = \{12, 10, 20, 34, 38, 40\}$
 - $K = 2$; Initially $\mu_0 = 10, \mu_1 = 20$
- $\text{Map}(X, \mu_0, \mu_1) \rightarrow (0, 12), (0, 10), (1, 20), (1, 34), (1, 38), (1, 40)$

- $\text{Shuffle}(\) \rightarrow L_0 = (0; 12, 10), L_1 = (1; 20, 34, 38, 40)$ (automatically done by the spark framework)
- $\text{Reduce}(L_0, L_1) \rightarrow \mu_0 = 11, \mu_1 = 32$ (Using ReduceByKey operation)
- $\text{Map}(X, \mu_0, \mu_1) \rightarrow (0, 12), (0, 10), (0, 20), (1, 34), (1, 38), (1, 40)$
- ...

Template for k-means program



```
import sys
from pyspark import SparkContext

def mapToCluster(data, means):
    #data -> a single integer value.
    #means -> list of the mean values.
    #return the mean value to which this data point belongs to
    return 0.0

def updatemeans(data1, data2):
    #data1,data2 -> tuple of format (meanvalue, count)
    #give (avg1, n1), (avg2, n2), new average will be (n1*avg1 + n2*avg2)/(n1+n2)
    return (newavg,newcount)

if __name__ == "__main__":
    if len(sys.argv) != 3:
        print(str(len(sys.argv))+"Usage: kmeans <datafile> <initialmeanfile>")
        exit(-1)

    #Create a sparkcontext
    sc = SparkContext(appName="kmeans")
    #load data from the text file
    data = sc.textFile(sys.argv[1]).cache()
    #load initial mean values from the text file
    means = sc.textFile(sys.argv[2])
    #We cannot directory use RDD. It should first be converted into a list to be iterated upon.
    meansList = means.collect()

    #we will run 50 iterations for calculating k means.
    numiter = 50


    for i in range(numiter):
        #For each data point create a tuple of the format (meanvalue, (datapoint, 1))
        clustermap = data.map(lambda p: (mapToCluster(p,meansList), (p,1)))
        #Use reduce operation to calculate new mean value for all the datapoint belonging to the same key
        newmeans = clustermap.reduceByKey(updatemeans)
        #Create a list from the RDD
        meansTupleList = newmeans.collect()
        meansList = []
        for mi in meansTupleList:
            meansList.append(mi[1][0])

    finalclustermap = data.map(lambda p: (mapToCluster(p,meansList), p)).sortByKey()
    finalclustermap.saveAsTextFile("output");
```




Triangle Counting (1)

- Input: directed graph
- Output: how many graph triangles each vertex belongs to






Triangle Counting (2)

- Map

- $X: \{x_1, x_2, \dots, x_N\}$ be a list relating vertex X with vertices $x_1, x_2 \dots x_n$

- Produce key-value pairs (k, v)

- k : neighbour of X
 - v : x_i
 - E.g. 3:3 \rightarrow map $\rightarrow (0,3), (2,3)$
 - Number of key-value pairs produced for X : $|d^+(X)| \times N$, where d^+ is the out-degree of X




Produced
pairs(*,3):
(0,3) (2,3)



Triangle Counting (3)


- Reduce
 - Collect all the key-value pairs produced in the previous Map step and produce $L_k = (k; v_1, v_2, \dots)$
 - Eg.
 - $L_0 = (0; 2, 3)$
 - $L_1 = (1; 0)$
 - $L_2 = (2; 1, 3)$
 - $L_3 = \dots$





Triangle Counting (4)

- The algorithm has 3 rounds
- Each round runs a map reduce step as described previously
- Interpretation of the rounds:
 - Input to Round i : $(X; \{x_1, x_2, \dots, x_N\})$: $\{x_1, x_2, \dots, x_N\}$ denote all the vertices from which we can reach X in $i - 1$ steps.
 - Map: (d, s) , all pairs of vertices such that we can reach d from s in i steps.
 - Reduce: $(X; \{x_1, x_2, \dots, x_N\})$: $\{x_1, x_2, \dots, x_N\}$ denote all the vertices from which we can reach X in i steps. Used as input for map in the next round.






Triangle Counting (5)

- Input: 0-1, 1-2, 1-3, 2-0, 3-0, 3-2
- Round 1:
 - $X_n : (0;0), (1;1), (2;2), (3;3)$
 - Map: (1,0), (2,1), (3,1), (0,2), (0,3), (2,3)
 - Reduce: (1;0), (2;(1,3)), (0;(3,2)), (3;1)
- Round 2:
 - $X_n : (1;0), (2;(1,3)), (0;(3,2)), (3;1)$
 - Map: (2,0), (3,0), (0,1), (0,3), (0,1), (2,1), (1,2), (1,3)
 - Reduce: (2;(0,1)), (3;0), (0;(1,1,3)), (1;(2,3))

Map: Given
(a;b) produce
(destn(a),b)



(U,v):
U is the 2-hop
Neighbor of v.



Triangle Counting (6)

- Round 3:
 - $X_n : (2;(0,1)), (3;0),(0;(1,1,3)),(1;(2,3))$
 - Map: $(0,0), (0,1), (0,0), (2,0), (1,1), (1,3), (1,1), (2,2), (2,3), (3,2), (3,3)$
 - Reduce: $(0;(0,1,0)), (2;(0,2,3)), (1;(1,3,1)), (3;(2,3))$
- Output:
 - $0 \rightarrow 2$
 - $2 \rightarrow 1$
 - $1 \rightarrow 2$
 - $3 \rightarrow 1$



Questions?

Thank you

Reference:

<https://www.cs.rutgers.edu/~pxk/417/notes/content/mapreduce.html>

<http://www.slideshare.net/mcorrea11/map-reduce-5584234>

<http://static.googleusercontent.com/media/research.google.com/en/us/archive/mapreduce-osdi04.pdf>

<https://github.com/himank/K-Means>